
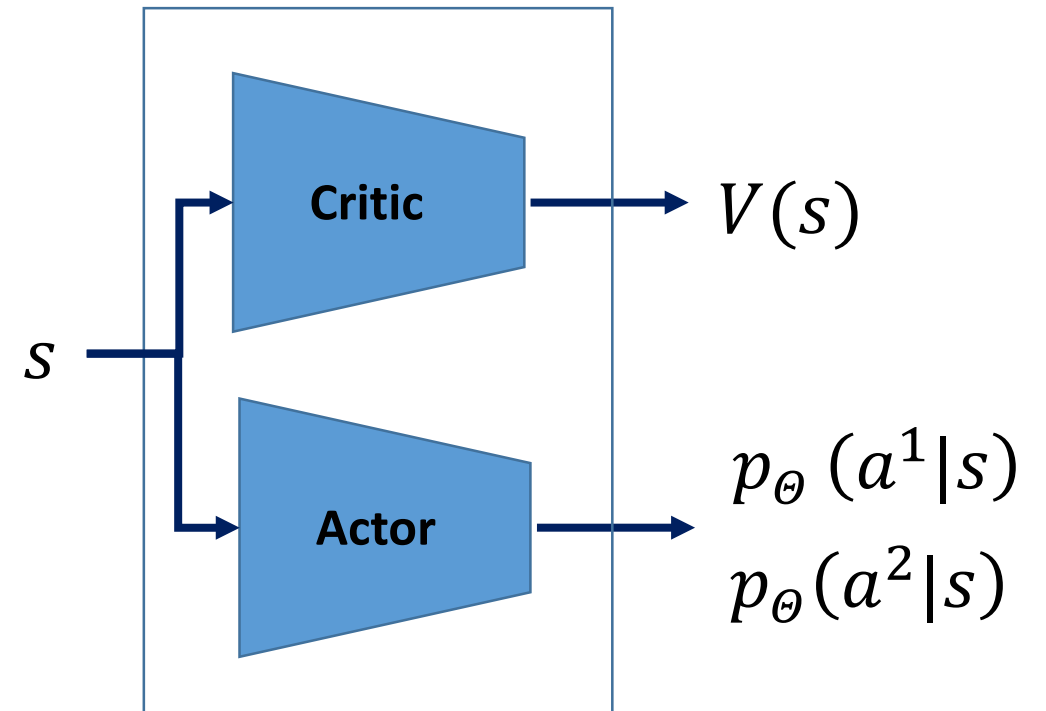
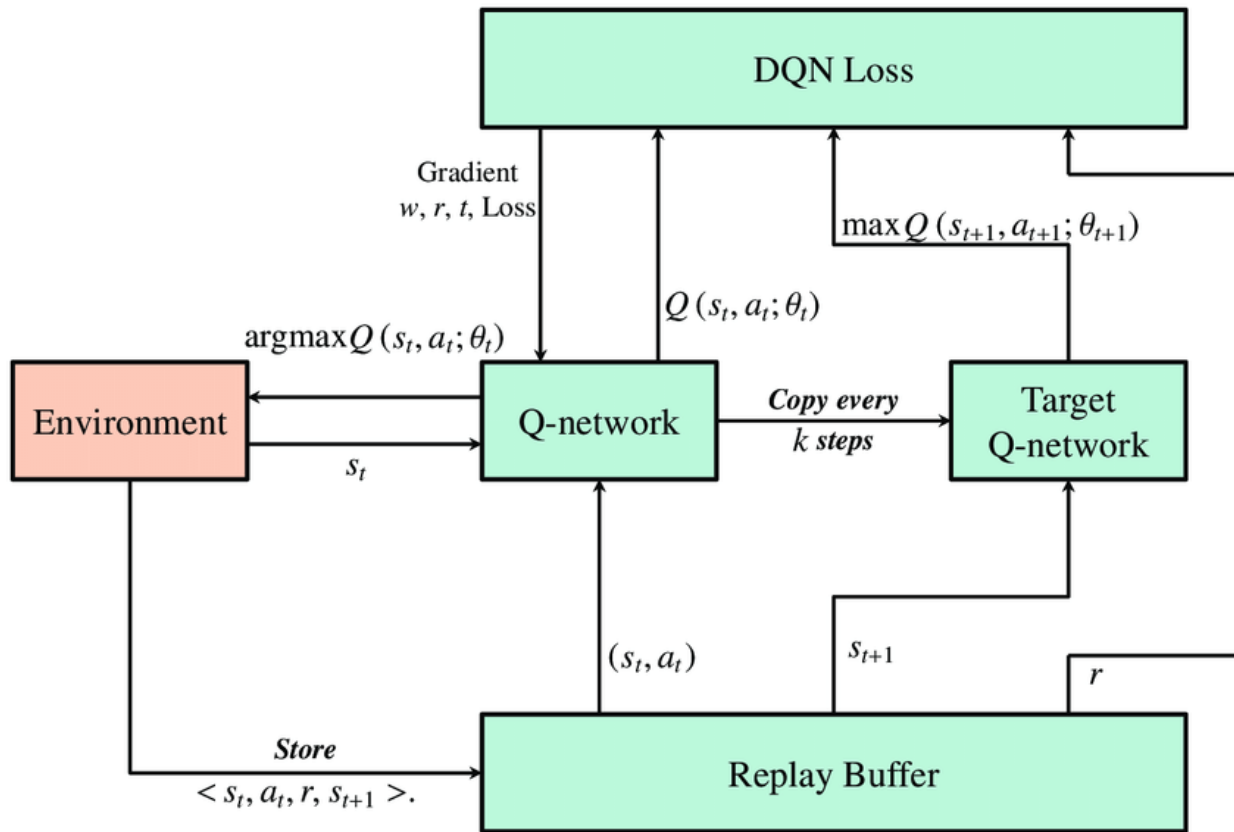


# Deep deterministic policy gradient



# Comparison – DQN and PPO-AC



# Introduction


Deep Deterministic Policy Gradient (DDPG) is an algorithm which concurrently learns a Q-function and a policy. It uses off-policy data and the Bellman equation to learn the Q-function, and uses the Q-function to learn the policy.

$$Q^*(s, a) = \mathbb{E} \left[ r(s, a) + \gamma \max_{a'} Q^*(s', a') \right]$$

$$a^*(s) = \arg \max_a Q^*(s, a)$$

Reference: <https://spinningup.openai.com/en/latest/algorithms/ddpg.html#deep-deterministic-policy-gradient>

# The Q-learning side of DDPG – Critic Net



$$L(\phi, \mathcal{D}) = \mathbb{E}_{(s, a, r, s', a') \sim \mathcal{D}} \left[ \left( Q_\phi(s, a) - (r_s^a + \gamma(1 - d) \max_{a'} Q_{\phi_{target}}(s', a')) \right)^2 \right]$$

# Trick one – replay buffer $\mathcal{D}$

$$L(\phi, \mathcal{D}) = \mathbb{E}_{(s,a,r,s',a') \sim \mathcal{D}} \left[ \left( Q_{\phi}(s, a) - (r_s^a + \gamma(1-d) \max_{a'} Q_{\phi}(s', a')) \right)^2 \right]$$

In order for the algorithm to have stable behavior, the replay buffer should be large enough to contain a wide range of experiences, but it may not always be good to keep everything. If you only use the very-most recent data, you will overfit to that and things will break; if you use too much experience, you may slow down your learning.

9. DDPG\_NN\_and\_MemoryBuffer.ipynb

```
class ReplayBuffer:
    """Fixed-size buffer to store experiences"""

    def __init__(self, action_space):
        """Initialize a ReplayBuffer"""
```

# Trick two – Target Network $\phi_{targ}$

$$L(\phi, \mathcal{D}) = \mathbb{E}_{(s,a,r,s',a') \sim \mathcal{D}} \left[ \left( Q_{\phi}(s, a) - (r_s^a + \gamma(1-d) \max_{a'} Q_{\phi_{targ}}(s', a')) \right)^2 \right]$$

Eval network  $\phi$

Target network  $\phi_{targ}$

9. DDPG\_Agent.ipynb

```
self.critic_local = Critic(state_size, action_size, r,
self.critic_target = Critic(state_size, action_size,
self.critic_optimizer = optim.Adam(self.critic_local.
```

$$\phi_{targ} \leftarrow \rho \phi_{targ} + (1 - \rho) \phi$$

```
self.soft_update(self.critic_local, self.critic_target, TAU)
self.soft_update(self.actor_local, self.actor_target, TAU)
```

# DDPG: calculating Max over actions in the Target

Computing the maximum over actions in the target is a challenge in continuous action spaces. DDPG deals with this by using a **target policy network** to compute an action which approximately maximizes  $Q_{\phi_{targ}}$ . The target policy network is found the same way as the target Q-function: by polyak averaging the policy parameters over the course of training.

$$L(\phi, \mathcal{D}) = \mathbb{E}_{(s,a,r,s',a') \sim \mathcal{D}} \left[ \left( Q_{\phi}(s, a) - (r_s^a + \gamma(1 - d)Q_{\phi_{targ}}(s', \mu_{\theta_{targ}}(s'))) \right)^2 \right]$$

$\mu_{\theta_{targ}}$ : target policy network

# DDPG: calculating Max over actions in the Target

$$L(\phi, \mathcal{D}) = \mathbb{E}_{(s,a,r,s',a') \sim \mathcal{D}} \left[ \left( Q_{\phi}(s, a) - (r_s^a + \gamma(1-d)Q_{\phi_{targ}}(s', \mu_{\theta_{targ}}(s'))) \right)^2 \right]$$

states, actions, rewards, next\_states, dones = experiences

actions\_next = self.actor\_target(next\_states)  $\mu_{\theta_{targ}}(s')$

Q\_targets\_next = self.critic\_target(next\_states, actions\_next)  $Q_{\phi_{targ}}(s', \mu_{\theta_{targ}}(s'))$

Q\_targets = rewards + (GAMMA \* Q\_targets\_next \* (1 - dones))  $(r_s^a + \gamma(1-d)Q_{\phi_{targ}}(s', \mu_{\theta_{targ}}(s'))$

Q\_expected = self.critic\_local(states, actions)  $Q_{\phi}(s, a)$

critic\_loss = F.mse\_loss(Q\_expected, Q\_targets)

self.critic\_optimizer.zero\_grad()

critic\_loss.backward()

9. DDPG\_Agent.ipynb



# The policy learning side of DDPG

$$\max_{\theta} \mathbb{E}_{s \sim \mathcal{D}} [Q_{\phi}(s, \mu_{\theta}(s))]$$

```
actions_pred = self.actor_local(states)            $\mu_{\theta}(s)$   
actor_loss = -self.critic_local(states, actions_pred).mean()   $Q_{\phi}(s, \mu_{\theta}(s))$   
self.actor_optimizer.zero_grad()  
actor_loss.backward()  
self.actor_optimizer.step()
```

9. DDPG\_Agent.ipynb

# Exploration vs. Exploitation

Because the policy is deterministic, if the agent were to explore on-policy, in the beginning it would probably not try a wide enough variety of actions to find useful learning signals. To make DDPG policies explore better, we add noise to their actions at training time.

9. DDPG\_Agent.ipynb

```
#adding noise for exploration!  
if add_noise:  
    acts += self.noise.sample()
```

# DDPG algorithm

---

**Algorithm 1** Deep Deterministic Policy Gradient

---

- 1: Input: initial policy parameters  $\theta$ , Q-function parameters  $\phi$ , empty replay buffer  $\mathcal{D}$
- 2: Set target parameters equal to main parameters  $\theta_{\text{targ}} \leftarrow \theta$ ,  $\phi_{\text{targ}} \leftarrow \phi$
- 3: **repeat**
- 4:   Observe state  $s$  and select action  $a = \text{clip}(\mu_{\theta}(s) + \epsilon, a_{\text{Low}}, a_{\text{High}})$ , where  $\epsilon \sim \mathcal{N}$
- 5:   Execute  $a$  in the environment
- 6:   Observe next state  $s'$ , reward  $r$ , and done signal  $d$  to indicate whether  $s'$  is terminal
- 7:   Store  $(s, a, r, s', d)$  in replay buffer  $\mathcal{D}$
- 8:   If  $s'$  is terminal, reset environment state.
- 9:   **if** it's time to update **then**
- 10:     **for** however many updates **do**
- 11:       Randomly sample a batch of transitions,  $B = \{(s, a, r, s', d)\}$  from  $\mathcal{D}$
- 12:       Compute targets

$$y(r, s', d) = r + \gamma(1 - d)Q_{\phi_{\text{targ}}}(s', \mu_{\theta_{\text{targ}}}(s'))$$

# DDPG algorithm

13:       Update Q-function by one step of gradient descent using

$$\nabla_{\phi} \frac{1}{|B|} \sum_{(s,a,r,s',d) \in B} (Q_{\phi}(s,a) - y(r,s',d))^2$$

14:       Update policy by one step of gradient ascent using

$$\nabla_{\theta} \frac{1}{|B|} \sum_{s \in B} Q_{\phi}(s, \mu_{\theta}(s))$$

15:       Update target networks with

$$\begin{aligned}\phi_{\text{targ}} &\leftarrow \rho \phi_{\text{targ}} + (1 - \rho) \phi \\ \theta_{\text{targ}} &\leftarrow \rho \theta_{\text{targ}} + (1 - \rho) \theta\end{aligned}$$

16:       **end for**

17:       **end if**

18: **until** convergence

---