



Design a NN to play 3D ball balancing

2. NN with policy interacts with 3D Ball (MLAgent 10).ipynb

Θ : neural network weights and biases



Policy gradient

3. NN with policy interacts with 3D Ball to collect training data (MLAgent_10).ipynb

$$\nabla \bar{R}_\theta \approx \frac{1}{N} \sum_{n=1}^N \sum_{t=1}^{20} \left(\sum_{t'}^{20} \gamma^{t'-t} r_{t'}^n - b \right) \nabla \log p_\theta(a_t^n | s_t^n)$$

Δ = reward + expected accumulated reward

gae = Δ + accumulated gae

Return = gae + expected accumulated reward

$$\Delta_{19} = r_{19} + (\gamma * v_{20} * mask_{19} - v_{19})$$

$$gae_{19 \sim 20} = \Delta_{19} + \gamma * \tau * mask_{19} * gae_{20}$$

$$return_{19} = gae_{19 \sim 20} + v_{19}$$

...

$$\Delta_{20} = r_{20} + (\gamma * v_{21} * mask_{20} - v_{20})$$

$$gae_{20} = \Delta_{20} + \gamma * \tau * mask_{20} * gae_{initial}$$

$$return_{20} = gae_{20} + v_{20}$$

$$\Delta_1 = r_1 + (\gamma * v_2 * mask_1 - v_1)$$


$$gae_{1 \sim 20} = \Delta_1 + \gamma * \tau * mask_1 * gae_{2 \sim 20}$$

$$return_1 = gae_{1 \sim 20} + v_1$$

Sampling efficiency problem of policy gradient

$$\nabla \bar{R}_{\Theta} \approx \frac{1}{N} \sum_{n=1}^N \sum_{t=1}^{T_n} \left(\sum_{t'}^{T_n} \gamma^{t'-t} r_{t'}^n - b \right) \nabla \log p_{\Theta}(a_t^n | s_t^n)$$

$$\Theta' \leftarrow \Theta + \eta \nabla \bar{R}_{\Theta}$$




$$\nabla \bar{R}_{\Theta'} \approx \frac{1}{N} \sum_{n=1}^N \sum_{t=1}^{T_n} \left(\sum_{t'}^{T_n} \gamma^{t'-t} r_{t'}^n - b \right) \nabla \log p_{\Theta'}(a_t^n | s_t^n)$$

PPO

4. NN optimization with PPO (MLAgent_10).ipynb

5. PPO (MLAgent_10) .ipynb



$$\max_{\theta'} \text{PPO2}(\theta')$$

$$\text{PPO2}(\theta') =$$

$$\sum_{(s_t, a_t)} \min \left(\frac{p_{\theta'}(a_t|s_t)}{p_\theta(a_t|s_t)} A^\theta(s_t, a_t), \text{clip} \left(\frac{p_{\theta'}(a_t|s_t)}{p_\theta(a_t|s_t)}, 1 - \varepsilon, 1 + \varepsilon \right) A^\theta(s_t, a_t) \right)$$

Use expected value to reduce sampling variance

$V^{\pi_\theta}(s_t^n)$: Expected long-term return of the current state s under policy π .

$Q^{\pi_\theta}(s_t^n, a_t^n)$: Expected long-term return of the current state s , taking action a under policy π .

$$\nabla \bar{R}_\theta \approx \frac{1}{N} \sum_{n=1}^N \sum_{t=1}^{T_n} \left(\sum_{t'}^{T_n} \gamma^{t'-t} r_{t'}^n - b \right) \nabla \log p_\theta(a_t^n | s_t^n)$$

$V^{\pi_\theta}(s_t^n)$ Expected value of b

$E[G_t^n] = Q^{\pi_\theta}(s_t^n, a_t^n)$ Expected value of G_t^n

$G_t^n = \sum_{t'}^{T_n} \gamma^{t'-t} r_{t'}^n$ unstable when sampling amount is not large enough

We only need to estimate $V(s)$

$$\nabla \bar{R}_\theta \approx \frac{1}{N} \sum_{n=1}^N \sum_{t=1}^{T_n} \left(\underbrace{\sum_{t'}^{T_n} \gamma^{t'-t} r_{t'}^n}_{\substack{\downarrow \\ E[G_t^n] = Q^{\pi_\theta}(s_t^n, a_t^n)}} - \underbrace{b}_{\substack{\uparrow \\ V^{\pi_\theta}(s_t^n)}} \right) \nabla \log p_\theta(a_t^n | s_t^n)$$

$$Q^{\pi_\theta}(s_t^n, a_t^n) = E[r_t^n + V^{\pi_\theta}(s_{t+1}^n)] = r_t^n + V^{\pi_\theta}(s_{t+1}^n)$$


$$Q^{\pi_\theta}(s_t^n, a_t^n) - V^{\pi_\theta}(s_t^n) = r_t^n + V^{\pi_\theta}(s_{t+1}^n) - V^{\pi_\theta}(s_t^n)$$

$$A^\theta(s_t, a_t) = (r_t^n + V^{\pi_\theta}(s_{t+1}^n) - V^{\pi_\theta}(s_t^n))$$

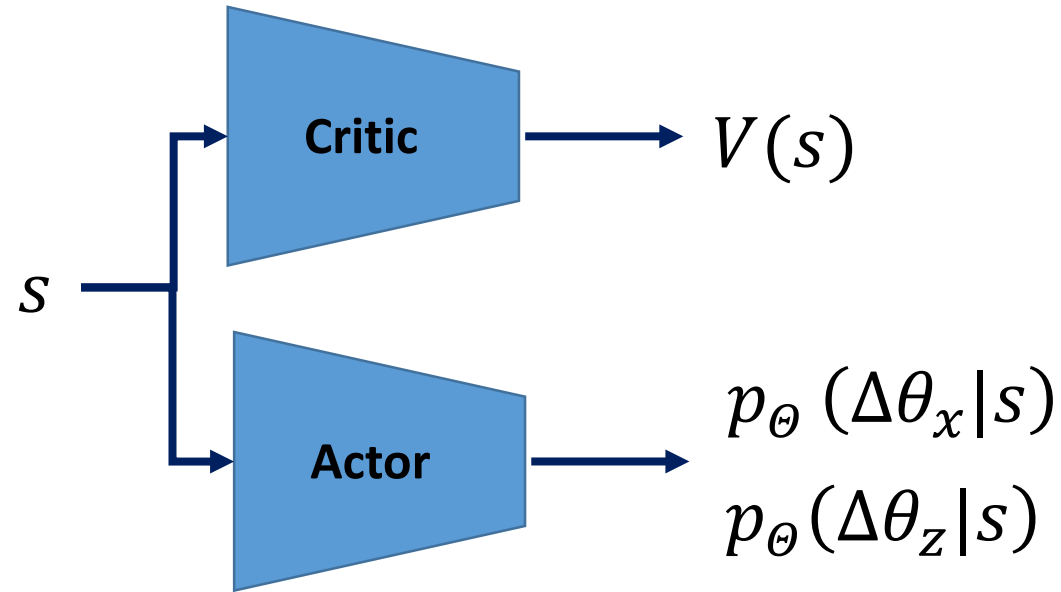
A2C (Advantage Actor Critic)

Actor – Learns the best actions (that can have maximum long-term rewards)

Critic – Learns the expected value of the long-term reward.



We need to know the true answer of $V(s)$



$$\nabla \bar{R}_{\theta} \approx \frac{1}{N} \sum_{n=1}^N \sum_{t=1}^{T_n} \left(\sum_{t'}^{T_n} \gamma^{t'-t} r_{t'}^n - b \right) \nabla \log p_{\theta}(a_t^n | s_t^n)$$

$$A^{\theta}(s_t, a_t) = (r_t^n + V^{\pi_{\theta}}(s_{t+1}^n) - V^{\pi_{\theta}}(s_t^n))$$

Use temporal difference to calculate $V(s)$

$$\nabla \bar{R}_\theta \approx \frac{1}{N} \sum_{n=1}^N \sum_{t=1}^{T_n} \left(\sum_{t'}^{T_n} \gamma^{t'-t} r_{t'}^n - b \right) \nabla \log p_\theta(a_t^n | s_t^n)$$

$$G_t^n = \sum_{t'}^{T_n} \gamma^{t'-t} r_{t'}^n$$

$$V^{\pi_\theta}(s_a) \leftrightarrow G_a$$

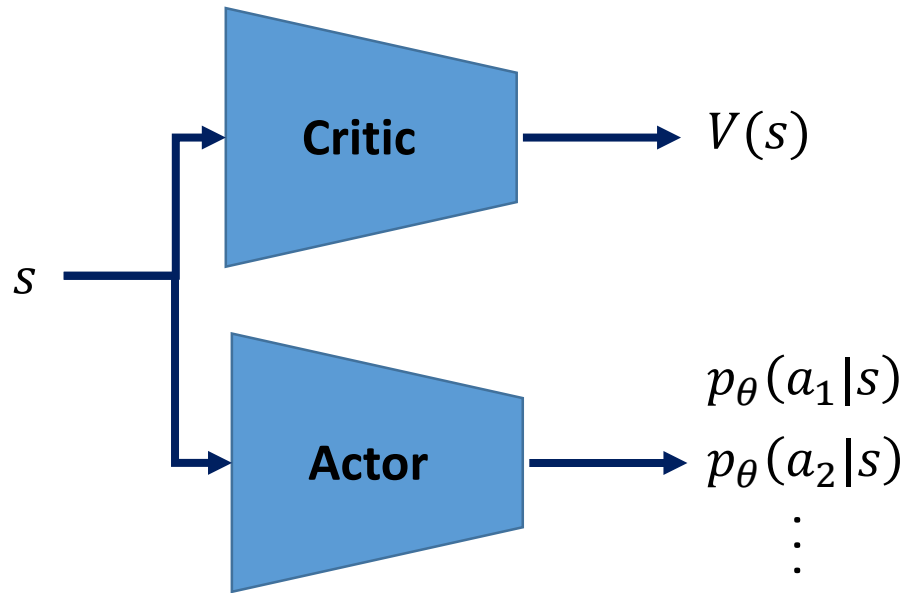
Monte-Carlo approach

$$V^{\pi_\theta}(s_t) + r_t = V^{\pi_\theta}(s_{t+1})$$

Temporal-difference approach

$$V^{\pi_\theta}(s_t) - V^{\pi_\theta}(s_{t+1}) \leftrightarrow r_t$$

Define loss function to optimize the Actor-Critic networks




$$L = L_{\pi} + c_v L_v$$

$$Loss_{\pi} = \sum_{(s_t, a_t)} \min \left(\frac{p_{\theta}(a_t|s_t)}{p_{\theta'}(a_t|s_t)} A^{\theta'}(s_t, a_t), \text{clip} \left(\frac{p_{\theta}(a_t|s_t)}{p_{\theta'}(a_t|s_t)}, 1 - \varepsilon, 1 + \varepsilon \right) A^{\theta'}(s_t, a_t) \right)$$

$$A^{\theta}(s_t, a_t) = r_t^n + \gamma V^{\pi_{\theta}}(s_{t+1}^n) - V^{\pi_{\theta}}(s_t^n)$$

$$Loss_v = \left(r_t^n + \gamma V^{\pi_{\theta}}(s_{t+1}^n) - V^{\pi_{\theta}}(s_t^n) \right)^2$$

Add entropy-based regularization



$$L = L_{\pi} + c_v L_v + c_{reg} L_{reg}$$

PyTorch implementation of A2C

6. A2C (MLAgent_10).ipynb

3DBall.yaml

behaviors:

3DBall:

trainer_type: ppo

hyperparameters:

batch_size: 64

buffer_size: 12000

learning_rate: 0.0003

beta: 0.001 c_{reg}

epsilon: 0.2 ϵ

lambda: 0.99 τ

num_epoch: 3

learning_rate_schedule: linear

network_settings:

normalize: true

hidden_units: 128

num_layers: 2

vis_encode_type: simple

reward_signals:

extrinsic:

gamma: 0.99 γ

strength: 1.0

keep_checkpoints: 5

max_steps: 50000

time_horizon: 1000

summary_freq: 5000

threaded: true

N_STATES = 8

N_ACTIONS = 2

N_AGENTS = 3 *#My Unity*

BATCH_SIZE = 64

BUFFER_SIZE = 12000

LEARNING_RATE = 0.0003

BETA = 0.001

EPSILON = 0.2

LAMBDA = 0.99

N_EPOCH = 3

GAMMA = 0.99

MAX_STEPS = 5000 *#50000*

TIME_HORIZON = 50 *#1000*

PyTorch implementation of A2C


```
class Net(nn.Module):
    def __init__(self, ):
        super(Net, self).__init__()

        self.critic = nn.Sequential(
            nn.Linear(N_STATES, 128),
            nn.LayerNorm(128),
            nn.Linear(128, 128),
            nn.LayerNorm(128),
            nn.Linear(128, 1)
        )

        self.actor = nn.Sequential(
            nn.Linear(N_STATES, 128),
            nn.LayerNorm(128),
            nn.Linear(128, 128),
            nn.LayerNorm(128),
            nn.Linear(128, N_ACTIONS)
        )

        self.log_std = nn.Parameter(torch.ones(1,
        self.apply(init_weights)

    def forward(self, x):
        value = self.critic(x)
        mu     = self.actor(x)
        std    = self.log_std.exp().expand_as(mu)
        dist   = Normal(mu, std)
        return dist, value
```



PyTorch implementation of A2C

```

dist, value = net(batch_state.to(device))
critic_loss = (batch_return.to(device) - value).pow(2).mean()
entropy = dist.entropy().mean()
batch_action = dist.sample()
batch_new_log_probs = dist.log_prob(batch_action)
ratio = (batch_new_log_probs - batch_old_log_probs.to(device)).exp()
surr1 = ratio * batch_advantage.to(device)
surr2 = torch.clamp(ratio, 1.0 - clip_param, 1.0 + clip_param) * batch_advantage
actor_loss = - torch.min(surr1, surr2).mean()
loss = 0.5 * critic_loss + actor_loss - 0.001 * entropy

```

$$L = L_{\pi} + c_v L_v + c_{reg} L_{reg}$$

$$Loss_v = (r_t^n + \gamma V^{\pi_{\theta}}(s_{t+1}^n) - V^{\pi_{\theta}}(s_t^n))^2$$

$$Loss_{\pi} = \sum_{(s_t, a_t)} \min \left(\frac{p_{\theta}(a_t|s_t)}{p_{\theta'}(a_t|s_t)} A^{\theta'}(s_t, a_t), \text{clip} \left(\frac{p_{\theta}(a_t|s_t)}{p_{\theta'}(a_t|s_t)}, 1 - \varepsilon, 1 + \varepsilon \right) A^{\theta'}(s_t, a_t) \right)$$

PyTorch implementation of A2C

```
while (frame_idx < MAX_STEPS):
    print("\nframe idx = ", frame_idx)
    print("Interacts with Unity to collect training data")
    states, actions, log_probs, values, rewards, masks, next_state = collect_training_data (N_AGEN
    _, next_value = net(next_state.to(device))

    print("Compute GAE of these training data set")
    returns = compute_gae(TIME_HORIZON, next_value, rewards, masks, values, GAMMA, LAMBD)

    returns = torch.cat(returns).detach()
    log_probs = torch.cat(log_probs).detach()
    values = torch.cat(values).detach()
    states = torch.cat(states)
    actions = torch.cat(actions)
    advantages = returns - values

    print("Optimize NN with PPO")
    critic_loss, actor_loss = ppo_update(N_EPOCH, BATCH_SIZE, states, actions, log_probs, returns,
    CriticLossLst.append(critic_loss)
    ActorLossLst.append(actor_loss)

    frame_idx += TIME_HORIZON
```


Use Tensorboard to visualize loss

6.1. A2C use Tensorboard (MLAgent_10).ipynb

Use Tensorboard to visualize loss

```
!echo The current directory is %CD%
```

The current directory is C:\Users\ADMIN\Google 雲端硬碟\0

```
# Run python terminal window  
# Change directory to this ipython notebook directory  
# tensorboard --logdir=runs
```


```
# Web browser: localhost:6006
```

```
from torch.utils.tensorboard import SummaryWriter  
writer = SummaryWriter()
```


```
critic_loss, actor_loss = ppo_update(N_EPOCH, BATCH_SIZE,  
writer.add_scalar("Actor loss", actor_loss, frame_idx)  
writer.add_scalar("Critic loss", critic_loss, frame_idx)
```

Use Tensorboard to visualize loss

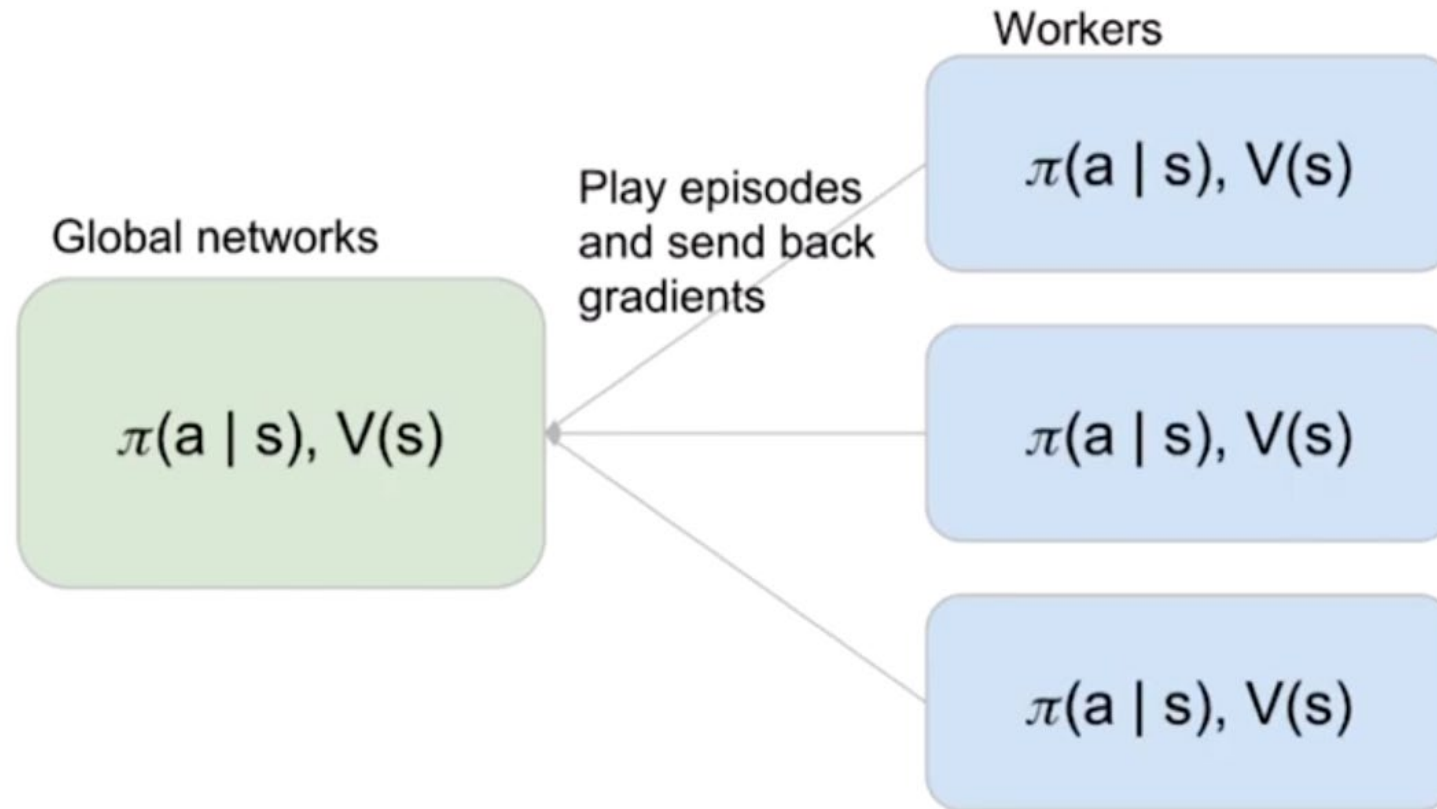
Actor loss



Critic loss



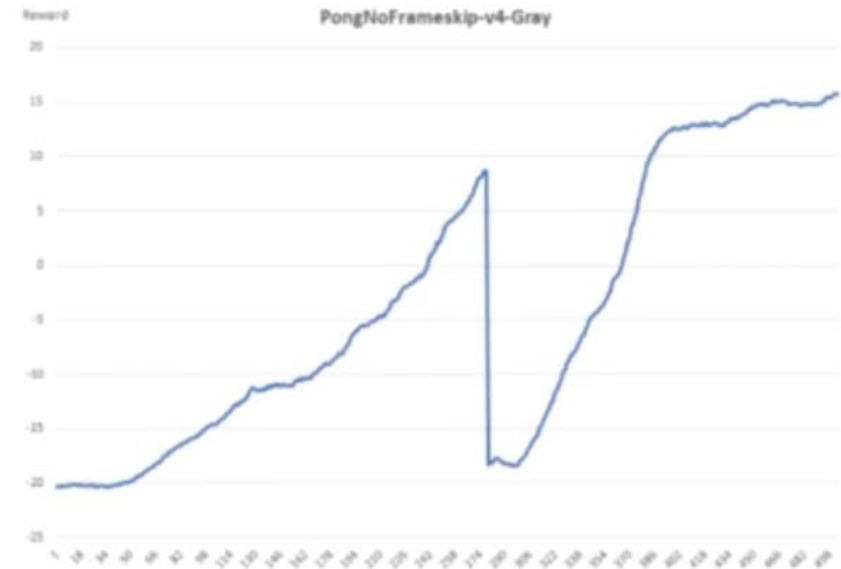
A3C (Asynchronous Advantage Actor Critic)



Reference: <https://youtu.be/iCV3vOl8IMk>

A3C

- Each episode will progress randomly
- Each action is sampled probabilistically
- Occasionally, performance of agent can drop off due to bad update
 - Well, this can still happen with A3C so don't think you are immune



Reference: <https://youtu.be/iCV3vOl8IMk>

A3C

- DQN is also interested in stabilizing learning
- Techniques:
 - Freezing target network
 - Experience replay buffer
- Use experience replay to look at multiple examples per training step
- A3C simply achieves stability using a different method (parallel agents)
- Both solve the problem: how to make neural networks work as function approximators in classic RL algorithms?

Reference: <https://youtu.be/iCV3vOl8IMk>

A3C

- Remember: the theory part is not new, just need to create multiple parallel agents and asynchronously update/copy parameters
- 3 files:
 - main.py (master file; global policy and value networks)
 - Create and coordinate workers
 - worker.py (contains local policy and value networks)
 - Copy weights from global nets
 - Play episodes
 - Send gradients back to master
 - nets.py
 - Definition of policy and value networks

A3C

main.py

Instantiate global policy and value networks

Check # CPUs available, create threads and workers

Initialize global thread-safe counter, so every worker knows when to quit (when # of total steps reaches a max.)

Reference: <https://youtu.be/iCV3vOl8IMk>

worker.py

```
def run():  
    in a loop:  
        copy params from global nets to local nets  
        run N steps of game (and store the data - s, a, r, s')  
        using gradients wrt local net, update the global net
```

Conceptually, it's like:

$$1) \quad g_{local} = \frac{\partial L(\theta_{local})}{\partial \theta_{local}}$$

$$2) \quad \theta_{global} = \theta_{global} - \eta g_{local}$$

But in reality, we'll use RMSprop

PPO-A3C

- PG $\nabla \bar{R}_\theta$
- Tips to reduce bias and variance in estimating $\nabla \bar{R}_\theta$
- Off-policy to improve efficiency of calculating $\nabla \bar{R}_\theta$
- Proximal policy optimization (PPO)
- Actor-critic strategy to reduce sampling variance in calculating $\nabla \bar{R}_\theta$
- Use temporal difference to calculate $V(s)$
- A3C